

Multiple Linear Regression for Baseball Win Prediction

Homework #1

Naomi Buell Richie Rivera Alexander Simon Nana Frimpong

2026-03-01

1 Introduction

We analyze and model data on a professional baseball team from 1871 to 2006. Each record has the performance of the team for the given year, with all of the statistics adjusted to match the performance of a 162 game season. We build a multiple linear regression model on the training data to predict the number of wins for the team.

2 Methods

2.1 Data Exploration

The monneyball training data set contains 2276 rows and 17 features, and is 91.01 percent complete. The data are all numeric, and the target variable is `TARGET_WINS`, which represents the number of games won by the team in a season. The other variables represent various performance statistics for the team. See Table 1 for descriptive characteristics of the training data. Variables range from 0% to 92% missing. We address missingness by dropping columns and imputing in the Data Preparation section.

Note

NB: Not sure if `target_wins` represents the number of games won per season. Is each row a season? Is this the correct interpretation of the target variable?

Table 1: Summary statistics for all variables in the training dataset

Variable	Completeness rate	Mean	Min	P25	P50	P75	Max
INDEX	1.000	1300	1	630	1300	1900	2500

Table 1: Summary statistics for all variables in the training dataset

Variable	Completeness rate	Mean	Min	P25	P50	P75	Max
TARGET_WINS	1.000	81	0	71	82	92	150
TEAM_BAT- TING_H	1.000	1500	890	1400	1500	1500	2600
TEAM_BAT- TING_2B	1.000	240	69	210	240	270	460
TEAM_BAT- TING_3B	1.000	55	0	34	47	72	220
TEAM_BAT- TING_HR	1.000	100	0	42	100	150	260
TEAM_BAT- TING_BB	1.000	500	0	450	510	580	880
TEAM_BAT- TING_SO	0.960	740	0	550	750	930	1400
TEAM_BASERUN_SB	0.940	120	0	66	100	160	700
TEAM_BASERUN_CS	0.660	53	0	38	49	62	200
TEAM_BAT- TING_HBP	0.084	59	29	50	58	67	95
TEAM_PITCH- ING_H	1.000	1800	1100	1400	1500	1700	30000
TEAM_PITCH- ING_HR	1.000	110	0	50	110	150	340
TEAM_PITCH- ING_BB	1.000	550	0	480	540	610	3600
TEAM_PITCH- ING_SO	0.960	820	0	620	810	970	19000
TEAM_FIELD- ING_E	1.000	250	65	130	160	250	1900
TEAM_FIELD- ING_DP	0.870	150	52	130	150	160	230

The correlation between all variables are shown Figure 1. The most correlated variables with the target variable `TARGET_WINS` are `TEAM_BATTING_H`, `TEAM_BATTING_2B`, and `TEAM_BATTING_BB`, which may be strong predictors of wins. The most autocorrelated variables are `TEAM_BATTING_HR` and `TEAM_PITCHING_HR`, which makes sense as teams that hit many homeruns may also allow many homeruns, for example.




Figure 1: Correlation matrix of all variables in the training dataset

Note

NB: Consider adding Bar Chart or Box Plot of the data or anything else that you think is interesting.

2.2 Data Preparation

We handled high missingness in the variables `TEAM_BATting_HBP`, `TEAM_BASERUN_CS`, and `TEAM_FIELDING_DP` (91.6%, 33.9%, and 12.6% missing, respectively), which are not very correlated with the target variable, by excluding these variables. The remaining variables are less than 10% missing, so we will impute the missing values with the median of each variable. We avoid dropping `TEAM_BASERUN_SB` as it is the most correlated variable with the target variable that has missingness, and instead impute it to preserve this potentially strong predictor of wins.

The median imputation was picked based on the distributions of missing variables: `TEAM_BASERUN_SB` and `TEAM_PITCHING_SO` have a long right tail so the median is a better imputation method than the mean. Both variables have observed have extreme outliers pulling the mean towards a right direction compared to where most the data sits. When the distribution of variables reveals two distinct peaks as observed in `TEAM_BATting_SO`, we can refer to this distribution as bimodal. Since the mean of bimodal falls between two peaks, this area is observed to be the underrepresented region of the missing data. The median imputation is utilized to approach one of the two clusters of missing variables with the largest distribution of missign variables. Below are the histograms

for the distributions of these variables.




Figure 2: Histograms of variables with missingness selected for imputation

To begin by remove variables with high missingness variables and low correlation with `TARGET_WINS`: `TEAM_BATTING_HBP` (91.6% missing), `TEAM_BASERUN_CS` (33.9% missing), and `TEAM_FIELDING_DP` (12.6% missing).

```
# Drop high missing variables and low-correlation variables with the data
# these variables lack a solid relationship with response variable TARGET_WINS
moneyball_clean <- moneyball_train |>
  select(-TEAM_BATTING_HBP, -TEAM_BASERUN_CS, -TEAM_FIELDING_DP)
```

Before our predictor variables `TEAM_BASERUN_CS`, `TEAM_BASERUN_SB` and `TEAM_PITCHING_SO` are imputed, a binary indicator is created to preserve information that value was missing in first place. Our indicator for each flag follows (1 = missing 0 = present)

```
moneyball_clean <- moneyball_clean |>
  mutate(
    TEAM_BASERUN_SB_MISSING = as.integer(is.na(TEAM_BASERUN_SB)),
    TEAM_BATTING_SO_MISSING = as.integer(is.na(TEAM_BATTING_SO)),
    TEAM_PITCHING_SO_MISSING = as.integer(is.na(TEAM_PITCHING_SO))
  )
```

Imputation now occurs based on the distribution of missing data of the following variables `TEAM_BASERUN_SB`, `TEAM_PITCHING_SO`, and `TEAM_BATTING_SO`. `TEAM_BASERUN_SB`

and `TEAM_PITCHING_SO` are right skewed therefore it is imputed around the mean and `TEAM_BATTING_SO` is bimodal, so we can impute around with the media to avoid placing the imputed values in region between two peak where only few observation exist.

```
moneyball_clean <- moneyball_clean %>%
  mutate(
    # Mean imputation for right-tailed variables
    TEAM_BASERUN_SB = ifelse(
      is.na(TEAM_BASERUN_SB),
      mean(TEAM_BASERUN_SB, na.rm = TRUE),
      TEAM_BASERUN_SB
    ),
    TEAM_PITCHING_SO = ifelse(
      is.na(TEAM_PITCHING_SO),
      mean(TEAM_PITCHING_SO, na.rm = TRUE),
      TEAM_PITCHING_SO
    ),

    # Median imputation for bimodal variable distribution
    TEAM_BATTING_SO = ifelse(
      is.na(TEAM_BATTING_SO),
      median(TEAM_BATTING_SO, na.rm = TRUE),
      TEAM_BATTING_SO
    )
  )
```

When a distribution is observed to be skewed to the right, meaning a small number of outliers has influence on the response variable `TARGET_WINS`, the log transformation can be applied to balance the scale of distribution. Utilizing the log compress the distribution of missing variable closer to symmetric shape. This makes the relationship with `TARGET_WINS` more linear and easier for the model to capture. `log1p` is used to handle any zero without producing undefined results

```
# Log transform right-skewed variables
moneyball_clean <- moneyball_clean %>%
  mutate(
    LOG_TEAM_BASERUN_SB = log1p(TEAM_BASERUN_SB),
    LOG_TEAM_PITCHING_SO = log1p(TEAM_PITCHING_SO),
    LOG_TEAM_FIELDING_E = log1p(TEAM_FIELDING_E)
  )
```

We can create three new variables that captures negative and positive impact on the wins for the team.

```
# HR_DIFFERENTIAL and BAT_BB_SO_RATIO
moneyball_clean <- moneyball_clean |>
  mutate(
    # Net home run differential (offensive power vs pitching vulnerability)
    # The team that hits more home runs than it gives up will produce positive value
    # the team that allows more home runs than it hits will produce a negative
    # value, meaning they lost the game
    # Zero indicates the team is even in home runs
    HR_DIFFERENTIAL = TEAM_BATTING_HR - TEAM_PITCHING_HR,
    # Here we can measure how often a batter can judge which pitch they would take a swing at
    # If the ratio is greater than 1, it reveals the team walks more than they strike
    # showing a more cohesive and patient team
    # This may happen if a batter forces the pitchers to throw more pitches

    BAT_BB_SO_RATIO = ifelse(TEAM_BATTING_SO > 0,
                              TEAM_BATTING_BB / TEAM_BATTING_SO, NA)
  )

# Impute missing values in BAT_BB_SO_RATIO with median
moneyball_clean <- moneyball_clean |>
  mutate(
    BAT_BB_SO_RATIO = ifelse(is.na(BAT_BB_SO_RATIO),
                              median(BAT_BB_SO_RATIO, na.rm = TRUE), BAT_BB_SO_RATIO)
  )
```

Finally we can verify that there are no missing variables that remain based on the results of our imputation

```
# Verify no remaining missingness
moneyball_clean |>
  summarise(across(everything(), ~ sum(is.na(.)))) |>
  tidyr::pivot_longer(everything(),
                      names_to = "Variable",
                      values_to = "Missing") |>
  filter(Missing > 0) |>
  knitr::kable()
```

Table 2: Remaining missing values after imputation

Variable	Missing
----------	---------

Describe how you have transformed the data by changing the original variables or creating new variables. If you did transform the data or create new variables, discuss why you did this. Here are some possible transformations.

- b. *Create flags to suggest if a variable was missing*
- c. *Transform data by putting it into buckets*
- d. *Mathematical transforms such as log or square root (or use Box-Cox)* e. *Combine variables (such as ratios or adding or multiplying) to create new variables*

2.3 Build Models

Using the training data set, build at least three different multiple linear regression models, using different variables (or the same variables with different transformations). Since we have not yet covered automated variable selection methods, you should select the variables manually (unless you previously learned Forward or Stepwise selection, etc.). Since you manually selected a variable for inclusion into the model or exclusion into the model, indicate why this was done. Discuss the coefficients in the models, do they make sense? For example, if a team hits a lot of Home Runs, it would be reasonably expected that such a team would win more games. However, if the coefficient is negative (suggesting that the team would lose more games), then that needs to be discussed. Are you keeping the model even though it is counter intuitive? Why? The boss needs to know.

2.4 Select Models

Decide on the criteria for selecting the best multiple linear regression model. Will you select a model with slightly worse performance if it makes more sense or is more parsimonious? Discuss why you selected your model. For the multiple linear regression model, will you use a metric such as Adjusted R^2 , RMSE, etc.? Be sure to explain how you can make inferences from the model, discuss multi-collinearity issues (if any), and discuss other relevant model output. Using the training data set, evaluate the multiple linear regression model based on (a) mean squared error, (b) R^2 , (c) F -statistic, and (d) residual plots. Make predictions using the evaluation data set.

3 Results

We predict that...

4 Conclusion

In conclusion...

5 Appendix: R Code

5.1 1. Data Exploration

First, we preview the data.

```
moneyball_train |> head()

# A tibble: 6 x 17
  INDEX TARGET_WINS TEAM_BATTING_H TEAM_BATTING_2B TEAM_BATTING_3B
  <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
1     1         39      1445         194         39
2     2         70      1339         219         22
3     3         86      1377         232         35
4     4         70      1387         209         38
5     5         82      1297         186         27
6     6         75      1279         200         36
# i 12 more variables: TEAM_BATTING_HR <dbl>, TEAM_BATTING_BB <dbl>,
#   TEAM_BATTING_SO <dbl>, TEAM_BASERUN_SB <dbl>, TEAM_BASERUN_CS <dbl>,
#   TEAM_BATTING_HBP <dbl>, TEAM_PITCHING_H <dbl>, TEAM_PITCHING_HR <dbl>,
#   TEAM_PITCHING_BB <dbl>, TEAM_PITCHING_SO <dbl>, TEAM_FIELDING_E <dbl>,
#   TEAM_FIELDING_DP <dbl>
```

The data are all numeric, and the target variable is `TARGET_WINS`, which represents the number of wins for the team in that season. The other variables represent various performance statistics for the team.

We also check summary statistics for all variables in our training data, including checking missingness.

```
moneyball_train |> skim()
```

Table 3: Data summary

Name	moneyball_train
------	-----------------

Number of rows	2276
Number of columns	17
Column type frequency:	
numeric	17
Group variables	None

Variable type: numeric

skim_vari- able	n_miss- ing	com- plete_rate	mean	sd	p0	p25	p50	p75	p100	hist
INDEX	0	1.00	1268.46	736.35	1	630.75	1270.5	1915.50	2535	
TAR- GET_WINS	0	1.00	80.79	15.75	0	71.00	82.0	92.00	146	
TEAM_BAT- TING_H	0	1.00	1469.27	144.59	891	1383.00	1454.0	1537.25	2554	
TEAM_BAT- TING_2B	0	1.00	241.25	46.80	69	208.00	238.0	273.00	458	
TEAM_BAT- TING_3B	0	1.00	55.25	27.94	0	34.00	47.0	72.00	223	
TEAM_BAT- TING_HR	0	1.00	99.61	60.55	0	42.00	102.0	147.00	264	
TEAM_BAT- TING_BB	0	1.00	501.56	122.67	0	451.00	512.0	580.00	878	
TEAM_BAT- TING_SO	102	0.96	735.61	248.53	0	548.00	750.0	930.00	1399	
TEAM_BASERUN3\$B	0	0.94	124.76	87.79	0	66.00	101.0	156.00	697	
TEAM_BASERUN7\$S	0	0.66	52.80	22.96	0	38.00	49.0	62.00	201	
TEAM_BAT- TING_HBP	2085	0.08	59.36	12.97	29	50.50	58.0	67.00	95	
TEAM_PITCH- ING_H	0	1.00	1779.21	1406.84	1137	1419.00	1518.0	1682.50	30132	
TEAM_PITCH- ING_HR	0	1.00	105.70	61.30	0	50.00	107.0	150.00	343	
TEAM_PITCH- ING_BB	0	1.00	553.01	166.36	0	476.00	536.5	611.00	3645	

skim_vari- able	n_miss- ing	com- plete_rate	mean	sd	p0	p25	p50	p75	p100	hist
TEAM_PITCH- ING_SO	102	0.96	817.73	553.09	0	615.00	813.5	968.00	19278	
TEAM_FIELD- ING_E	0	1.00	246.48	227.77	65	127.00	159.0	249.25	1898	
TEAM_FIELD- ING_DP	286	0.87	146.39	26.23	52	131.00	149.0	164.00	228	

```
nrows <- moneyball_train |> nrow()
ncols <- moneyball_train |> length()
perc_complete <- round(moneyball_train |> n_complete() / nrows / ncols * 100, 2)
```

This training dataset contains 2276 rows and 17 features, and is 91.01 percent complete. The variables with missingness are:

```
moneyball_train |>
  summarise(across(everything(), ~ sum(is.na(.)) / nrows * 100)) |>
  pivot_longer(everything(), names_to = "variable", values_to = "missing_percentage") |>
  arrange(desc(missing_percentage)) |>
  filter(missing_percentage > 0)
```


```
# A tibble: 6 x 2
  variable      missing_percentage
  <chr>          <dbl>
1 TEAM_BATTING_HBP      91.6
2 TEAM_BASERUN_CS      33.9
3 TEAM_FIELDING_DP     12.6
4 TEAM_BASERUN_SB       5.76
5 TEAM_BATTING_SO       4.48
6 TEAM_PITCHING_SO      4.48
```

This missingness will need to be addressed in the data preparation step.

We also visualize correlations between all variables.

```
moneyball_train_cor <- moneyball_train |>
  cor(use = "pairwise.complete.obs")

corrplot(moneyball_train_cor, method = "circle", type = "upper", tl.cex = 0.7, sig.level = 0.05)
```



The most correlated variables with the target variable TARGET_WINS are:

```
moneyball_train_cor |>
  as.data.frame() |>
  rownames_to_column(var = "variable") |>
  select(variable, TARGET_WINS) |>
  filter(variable != "TARGET_WINS") |>
  arrange(desc(abs(TARGET_WINS))) |>
  head(10)
```

	variable	TARGET_WINS
1	TEAM_BATTING_H	0.3887675
2	TEAM_BATTING_2B	0.2891036
3	TEAM_BATTING_BB	0.2325599
4	TEAM_PITCHING_HR	0.1890137
5	TEAM_FIELDING_E	-0.1764848
6	TEAM_BATTING_HR	0.1761532
7	TEAM_BATTING_3B	0.1426084
8	TEAM_BASERUN_SB	0.1351389
9	TEAM_PITCHING_BB	0.1241745
10	TEAM_PITCHING_H	-0.1099371

These may be strong predictors of wins.

The most autocorrelated variables are:

```

moneyball_train_cor |>
  as.data.frame() |>
  rownames_to_column(var = "variable1") |>
  pivot_longer(-variable1, names_to = "variable2", values_to = "correlation") |>
  filter(variable1 != variable2) |>
  distinct(across(c(correlation)), .keep_all = TRUE) |>
  arrange(desc(abs(correlation))) |>
  head(20)

```

A tibble: 20 x 3


	variable1	variable2	correlation
	<chr>	<chr>	<dbl>
1	TEAM_BATTING_HR	TEAM_PITCHING_HR	0.969
2	TEAM_BATTING_HR	TEAM_BATTING_SO	0.727
3	TEAM_BATTING_3B	TEAM_BATTING_SO	-0.670
4	TEAM_PITCHING_H	TEAM_FIELDING_E	0.668
5	TEAM_BATTING_SO	TEAM_PITCHING_HR	0.667
6	TEAM_BATTING_BB	TEAM_FIELDING_E	-0.656
7	TEAM_BASERUN_SB	TEAM_BASERUN_CS	0.655
8	TEAM_BATTING_3B	TEAM_BATTING_HR	-0.636
9	TEAM_BATTING_HR	TEAM_FIELDING_E	-0.587
10	TEAM_BATTING_SO	TEAM_FIELDING_E	-0.585
11	TEAM_BATTING_3B	TEAM_PITCHING_HR	-0.568
12	TEAM_BATTING_H	TEAM_BATTING_2B	0.563
13	TEAM_BATTING_3B	TEAM_BASERUN_SB	0.534
14	TEAM_BATTING_HR	TEAM_BATTING_BB	0.514
15	TEAM_BATTING_3B	TEAM_FIELDING_E	0.510
16	TEAM_BASERUN_SB	TEAM_FIELDING_E	0.510
17	TEAM_FIELDING_E	TEAM_FIELDING_DP	-0.498
18	TEAM_BASERUN_SB	TEAM_FIELDING_DP	-0.497
19	TEAM_PITCHING_HR	TEAM_FIELDING_E	-0.493
20	TEAM_BATTING_BB	TEAM_PITCHING_BB	0.489

It makes sense that homeruns by batters (TEAM_BATTING_HR) and homeruns allowed (TEAM_PITCHING_HR) are correlated, as teams that hit many homeruns may also allow many homeruns, for example. We may need to avoid including these pairs of variables in the same model to prevent multicollinearity.


We also visualize the distribution of all variables.

```
# Loop through all variables and create histograms
for (var in names(moneyball_train)) {
  p <- ggplot(moneyball_train, aes_string(x = var)) +
    geom_histogram() +
    labs(title = paste("Distribution of", var), x = var, y = "Frequency") +
    theme_minimal() +
    theme(panel.grid = element_blank())
  print(p)
}
```


Distribution of INDEX




Distribution of TARGET_WINS







Distribution of TEAM_BATTING_H




Distribution of TEAM_BATTING_2B










Distribution of TEAM_BASERUN_SB




Distribution of TEAM_BASERUN_CS









Distribution of TEAM_PITCHING_HR



Distribution of TEAM_PITCHING_BB







5.2 2. Data Preparation

Next, we will prepare the data for modeling.

First, we handle missing data. Since we saw in the EDA that `TEAM_BATTING_HBP`, `TEAM_BASERUN_CS`, and `TEAM_FIELDING_DP` were 91.6%, 33.9%, and 12.6% missing, respectively, and not very correlated with the target variable, so we will drop these variables. The remaining variables are less than 10% missing, so we will impute the missing values with the median of each variable. `TEAM_BASERUN_SB` is the most correlated variable with the target variable that has missingness, so we will be sure to impute this variable rather than drop it.

The median imputation was picked based on the distributions of the variables with missingness: `TEAM_BASERUN_SB` and `TEAM_PITCHING_SO` have a long right tail so the median is a better imputation method than the mean. `TEAM_BATTING_SO` is bimodal, so the median is also a better imputation method than the mean for this variable. See the histograms above for the distributions of these variables.

```
moneyball_train_prepared <- moneyball_train |>
  select(-TEAM_BATTING_HBP, -TEAM_BASERUN_CS, -TEAM_FIELDING_DP) |>
  mutate(across(everything(), ~ ifelse(is.na(.), median(., na.rm = TRUE), .)))
```

5.3 3. Build Models

5.4 4. Select Models

Export assigned predictions (the number of wins for the team) for the evaluation data set for deliverable.