

Multiple Linear Regression for Baseball Win Prediction

Homework #1

Naomi Buell

Richie Rivera

Alexander Simon

2026-03-01

1 Introduction

We analyze and model data on a professional baseball team from 1871 to 2006. Each record has the performance of the team for the given year, with all of the statistics adjusted to match the performance of a 162 game season. We build a multiple linear regression model on the training data to predict the number of wins for the team.

2 Methods

2.1 Data Exploration

The mondeyball training data set contains 2276 rows and 17 features, and is 91.01 percent complete. The data are all numeric, and the target variable is `TARGET_WINS`, which represents the number of games won by the team in a season. The other variables represent various performance statistics for the team. See Table 1 for descriptive characteristics of the training data. Variables range from 0% to 92% missing. We address missingness by dropping columns and imputing in the Data Preparation section.

i Note

NB: Not sure if `target_wins` represents the number of games won per season. Is each row a season? Is this the correct interpretation of the target variable?

Table 1: Summary statistics for all variables in the training dataset

Variable	Completeness rate	Mean	Min	P25	P50	P75	Max
INDEX	1.000	1300	1	630	1300	1900	2500
TARGET_WINS	1.000	81	0	71	82	92	150

Table 1: Summary statistics for all variables in the training dataset

Variable	Completeness rate	Mean	Min	P25	P50	P75	Max
TEAM_BATTING_H	1.000	1500	890	1400	1500	1500	2600
TEAM_BATTING_2B	1.000	240	69	210	240	270	460
TEAM_BATTING_3B	1.000	55	0	34	47	72	220
TEAM_BATTING_BB	1.000	100	0	42	100	150	260
TEAM_BATTING_HR	1.000	500	0	450	510	580	880
TEAM_BATTING_SO	0.960	740	0	550	750	930	1400
TEAM_BASERUN_SB	0.940	120	0	66	100	160	700
TEAM_BASERUN_CS	0.660	53	0	38	49	62	200
TEAM_BATTING_HBP	0.084	59	29	50	58	67	95
TEAM_PITCHING_H	1.000	1800	1100	1400	1500	1700	30000
TEAM_PITCHING_HR	1.000	110	0	50	110	150	340
TEAM_PITCHING_BB	1.000	550	0	480	540	610	3600
TEAM_PITCHING_SO	0.960	820	0	620	810	970	19000
TEAM_FIELDING_E	1.000	250	65	130	160	250	1900
TEAM_FIELDING_DP	0.870	150	52	130	150	160	230

The correlation between all variables are shown Figure 1. The most correlated variables with the target variable TARGET_WINS are TEAM_BATTING_H, TEAM_BATTING_2B, and TEAM_BATTING_BB, which may be strong predictors of wins. The most autocorrelated variables are TEAM_BATTING_HR and TEAM_PITCHING_HR, which makes sense as teams that hit many homeruns may also allow many homeruns, for example.

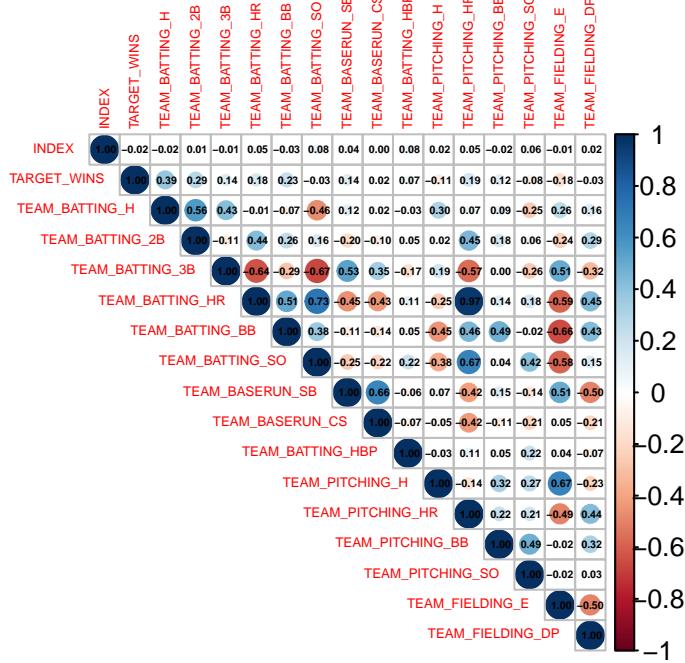


Figure 1: Correlation matrix of all variables in the training dataset

i Note

NB: Consider adding Bar Chart or Box Plot of the data or anything else that you think is interesting.

2.2 Data Preparation

We handled high missingness in the variables `TEAM_BATTING_HBP`, `TEAM_BASERUN_CS`, and `TEAM_FIELDING_DP` (91.6%, 33.9%, and 12.6% missing, respectively), which are not very correlated with the target variable, by excluding these variables. The remaining variables are less than 10% missing, so we will impute the missing values with the median of each variable. We avoid dropping `TEAM_BASERUN_SB` as it is the most correlated variable with the target variable that has missingness, and instead impute it to preserve this potentially strong predictor of wins.

The median imputation was picked based on the distributions of the variables with missingness: `TEAM_BASERUN_SB` and `TEAM_PITCHING_SO` have a long right tail so the median is a better imputation method than the mean. `TEAM_BATTING_SO` is bimodal, so the median is also a better imputation method than the mean for this variable. See the histograms above for the distributions of these variables.

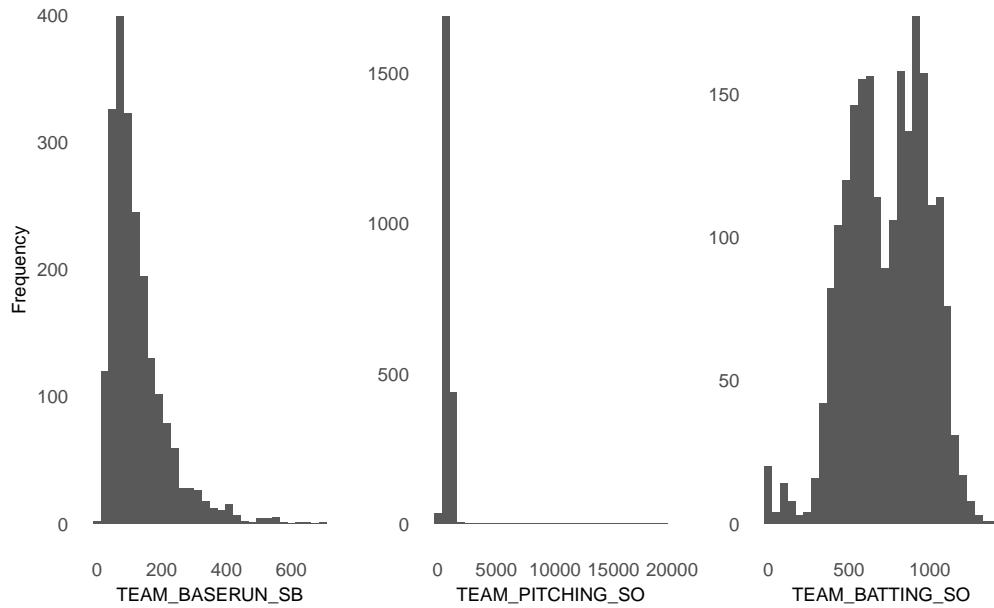


Figure 2: Histograms of variables with missingness selected for imputation

Describe how you have transformed the data by changing the original variables or creating new variables. If you did transform the data or create new variables, discuss why you did this. Here are some possible transformations.

- b. Create flags to suggest if a variable was missing
- c. Transform data by putting it into buckets
- d. Mathematical transforms such as log or square root (or use Box-Cox) e. Combine variables (such as ratios or adding or multiplying) to create new variables

2.3 Build Models

Using the training data set, build at least three different multiple linear regression models, using different variables (or the same variables with different transformations). Since we have not yet covered automated variable selection methods, you should select the variables manually (unless you previously learned Forward or Stepwise selection, etc.). Since you manually selected a variable for inclusion into the model or exclusion into the model, indicate why this was done. Discuss the coefficients in the models, do they make sense? For example, if a team hits a lot of Home Runs, it would be reasonably expected that such a team would win more games. However, if the coefficient is negative (suggesting that the team would lose more games), then that needs to be discussed. Are you keeping the model even though it is counter intuitive? Why? The boss needs to know.

2.4 Select Models

Decide on the criteria for selecting the best multiple linear regression model. Will you select a model with slightly worse performance if it makes more sense or is more parsimonious? Discuss why you selected your model. For the multiple linear regression model, will you use a metric such as Adjusted R², RMSE, etc.? Be sure to explain how you can make inferences from the model, discuss multi-collinearity issues (if any), and discuss other relevant model output. Using the training data set, evaluate the multiple linear regression model based on (a) mean squared error, (b) R², (c) F-statistic, and (d) residual plots. Make predictions using the evaluation data set.

3 Results

We predict that...

4 Conclusion

In conclusion...

5 Appendix: R Code

5.1 1. Data Exploration

First, we preview the data.

```
moneyball_train |> head()

# A tibble: 6 x 17
  INDEX TARGET_WINS TEAM_BATTING_H TEAM_BATTING_2B TEAM_BATTING_3B
  <dbl>      <dbl>       <dbl>       <dbl>       <dbl>
1     1         39        1445        194        39
2     2         70        1339        219        22
3     3         86        1377        232        35
4     4         70        1387        209        38
5     5         82        1297        186        27
6     6         75        1279        200        36
# i 12 more variables: TEAM_BATTING_HR <dbl>, TEAM_BATTING_BB <dbl>,
#   TEAM_BATTING_SO <dbl>, TEAM_BASERUN_SB <dbl>, TEAM_BASERUN_CS <dbl>,
#   TEAM_BATTING_HBP <dbl>, TEAM_PITCHING_H <dbl>, TEAM_PITCHING_HR <dbl>,
#   TEAM_PITCHING_BB <dbl>, TEAM_PITCHING_SO <dbl>, TEAM_FIELDING_E <dbl>,
#   TEAM_FIELDING_DP <dbl>
```

The data are all numeric, and the target variable is TARGET_WINS, which represents the number of

wins for the team in that season. The other variables represent various performance statistics for the team.

We also check summary statistics for all variables in our training data, including checking missingness.

```
moneyball_train |> skim()
```

Table 2: Data summary

Name	moneyball_train
Number of rows	2276
Number of columns	17
Column type frequency:	
numeric	17
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
INDEX	0	1.00	1268.46	736.35	1	630.75	1270.5	1915.50	2535	
TAR-	0	1.00	80.79	15.75	0	71.00	82.0	92.00	146	
GET_WINS										
TEAM_BATTING_H	0	1.00	1469.27	144.59	891	1383.00	1454.0	1537.25	2554	
TEAM_BATTING_2B	0	1.00	241.25	46.80	69	208.00	238.0	273.00	458	
TEAM_BATTING_3B	0	1.00	55.25	27.94	0	34.00	47.0	72.00	223	
TEAM_BATTING_HR	0	1.00	99.61	60.55	0	42.00	102.0	147.00	264	
TEAM_BATTING_BB	0	1.00	501.56	122.67	0	451.00	512.0	580.00	878	
TEAM_BATTING_SO	102	0.96	735.61	248.53	0	548.00	750.0	930.00	1399	
TEAM_BASERUN1SB	0.94	124.76	87.79	0	66.00	101.0	156.00	697		
TEAM_BASERUN7CS	0.66	52.80	22.96	0	38.00	49.0	62.00	201		
TEAM_BATTING_HBP	2085	0.08	59.36	12.97	29	50.50	58.0	67.00	95	

skim_vari- able	n_miss- ing	com- plete_rate	mean	sd	p0	p25	p50	p75	p100	hist
TEAM_PITCH- ING_H	0	1.00	1779.21	1406.84	1137	1419.00	1518.0	1682.50	30132	
TEAM_PITCH- ING_HR	0	1.00	105.70	61.30	0	50.00	107.0	150.00	343	
TEAM_PITCH- ING_BB	0	1.00	553.01	166.36	0	476.00	536.5	611.00	3645	
TEAM_PITCH- ING_SO	102	0.96	817.73	553.09	0	615.00	813.5	968.00	19278	
TEAM_FIELD- ING_E	0	1.00	246.48	227.77	65	127.00	159.0	249.25	1898	
TEAM_FIELD- ING_DP	286	0.87	146.39	26.23	52	131.00	149.0	164.00	228	

```

nrows <- moneyball_train |> nrow()
ncols <- moneyball_train |> length()
perc_complete <- round(moneyball_train |> n_complete() / nrows / ncols * 100, 2)

```

This training dataset contains 2276 rows and 17 features, and is 91.01 percent complete. The variables with missingness are:

```

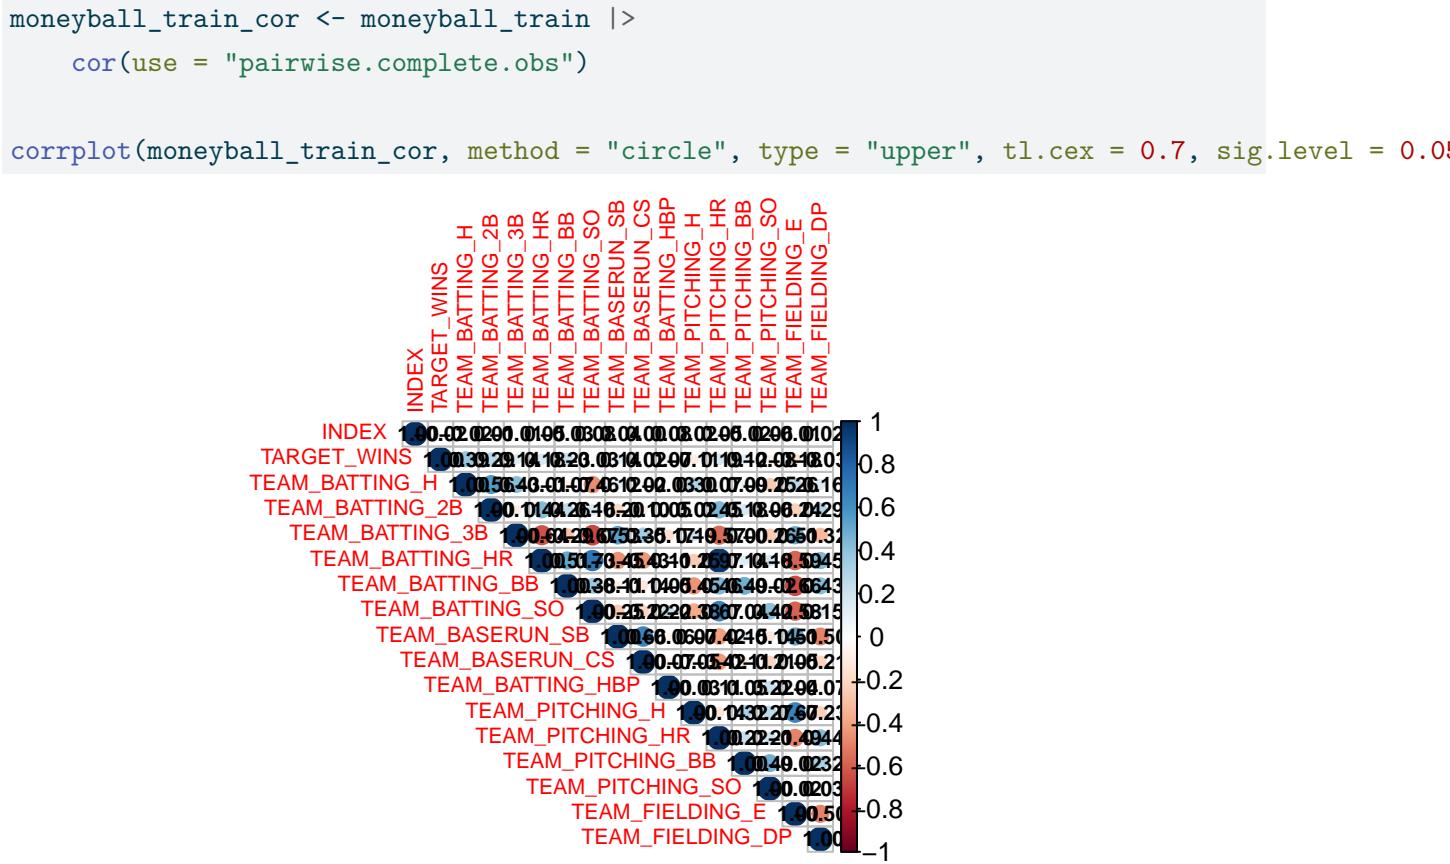
moneyball_train |>
  summarise(across(everything(), ~ sum(is.na(.)) / nrows * 100)) |>
  pivot_longer(everything(), names_to = "variable", values_to = "missing_percentage") |>
  arrange(desc(missing_percentage)) |>
  filter(missing_percentage > 0)

# A tibble: 6 x 2
  variable      missing_percentage
  <chr>                <dbl>
1 TEAM_BATTING_HBP        91.6
2 TEAM_BASERUN_CS         33.9
3 TEAM_FIELDING_DP        12.6
4 TEAM_BASERUN_SB         5.76
5 TEAM_BATTING_SO          4.48
6 TEAM_PITCHING_SO          4.48

```

This missingness will need to be addressed in the data preparation step.

We also visualize correlations between all variables.



The most correlated variables with the target variable TARGET_WINS are:

```
moneyball_train_cor |>
  as.data.frame() |>
  rownames_to_column(var = "variable") |>
  select(variable, TARGET_WINS) |>
  filter(variable != "TARGET_WINS") |>
  arrange(desc(abs(TARGET_WINS))) |>
  head(10)
```

	variable	TARGET_WINS
1	TEAM_BATTING_H	0.3887675
2	TEAM_BATTING_2B	0.2891036
3	TEAM_BATTING_BB	0.2325599
4	TEAM_PITCHING_HR	0.1890137
5	TEAM_FIELDING_E	-0.1764848
6	TEAM_BATTING_HR	0.1761532
7	TEAM_BATTING_3B	0.1426084
8	TEAM_BASERUN_SB	0.1351389
9	TEAM_PITCHING_BB	0.1241745

```
10 TEAM_PITCHING_H -0.1099371
```

These may be strong predictors of wins.

The most autocorrelated variables are:

```
moneyball_train_cor |>
  as.data.frame() |>
  rownames_to_column(var = "variable1") |>
  pivot_longer(-variable1, names_to = "variable2", values_to = "correlation") |>
  filter(variable1 != variable2) |>
  distinct(across(c(correlation)), .keep_all = TRUE) |>
  arrange(desc(abs(correlation))) |>
  head(20)
```

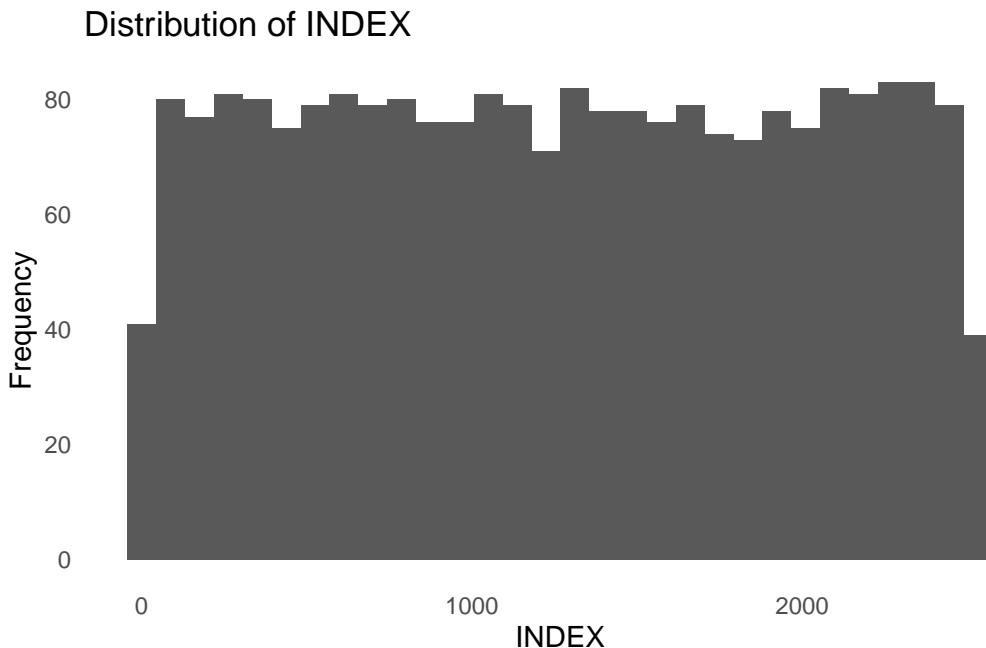
	variable1	variable2	correlation
	<chr>	<chr>	<dbl>
1	TEAM_BATTING_HR	TEAM_PITCHING_HR	0.969
2	TEAM_BATTING_HR	TEAM_BATTING_SO	0.727
3	TEAM_BATTING_3B	TEAM_BATTING_SO	-0.670
4	TEAM_PITCHING_H	TEAM_FIELDING_E	0.668
5	TEAM_BATTING_SO	TEAM_PITCHING_HR	0.667
6	TEAM_BATTING_BB	TEAM_FIELDING_E	-0.656
7	TEAM_BASERUN_SB	TEAM_BASERUN_CS	0.655
8	TEAM_BATTING_3B	TEAM_BATTING_HR	-0.636
9	TEAM_BATTING_HR	TEAM_FIELDING_E	-0.587
10	TEAM_BATTING_SO	TEAM_FIELDING_E	-0.585
11	TEAM_BATTING_3B	TEAM_PITCHING_HR	-0.568
12	TEAM_BATTING_H	TEAM_BATTING_2B	0.563
13	TEAM_BATTING_3B	TEAM_BASERUN_SB	0.534
14	TEAM_BATTING_HR	TEAM_BATTING_BB	0.514
15	TEAM_BATTING_3B	TEAM_FIELDING_E	0.510
16	TEAM_BASERUN_SB	TEAM_FIELDING_E	0.510
17	TEAM_FIELDING_E	TEAM_FIELDING_DP	-0.498
18	TEAM_BASERUN_SB	TEAM_FIELDING_DP	-0.497
19	TEAM_PITCHING_HR	TEAM_FIELDING_E	-0.493
20	TEAM_BATTING_BB	TEAM_PITCHING_BB	0.489

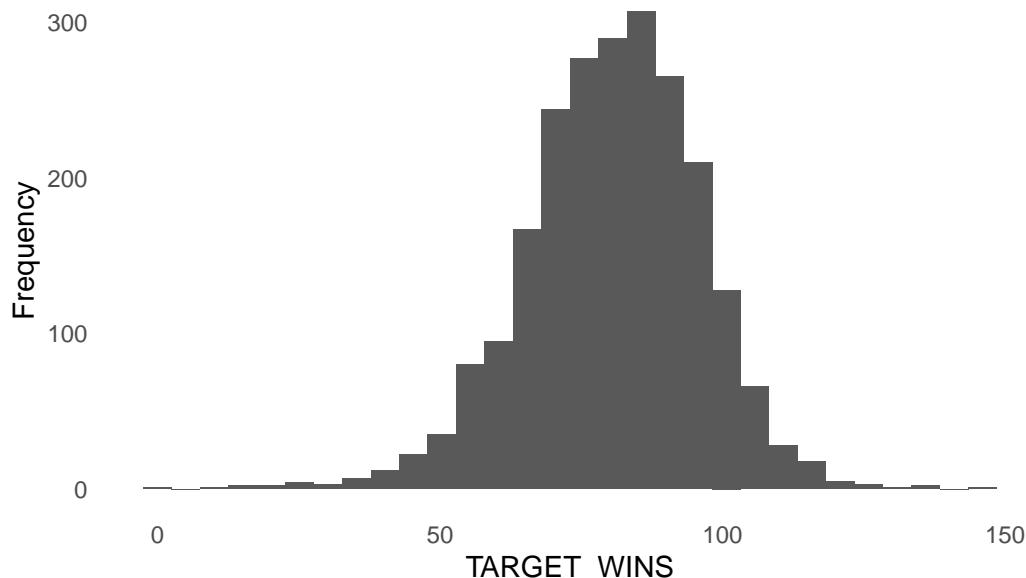
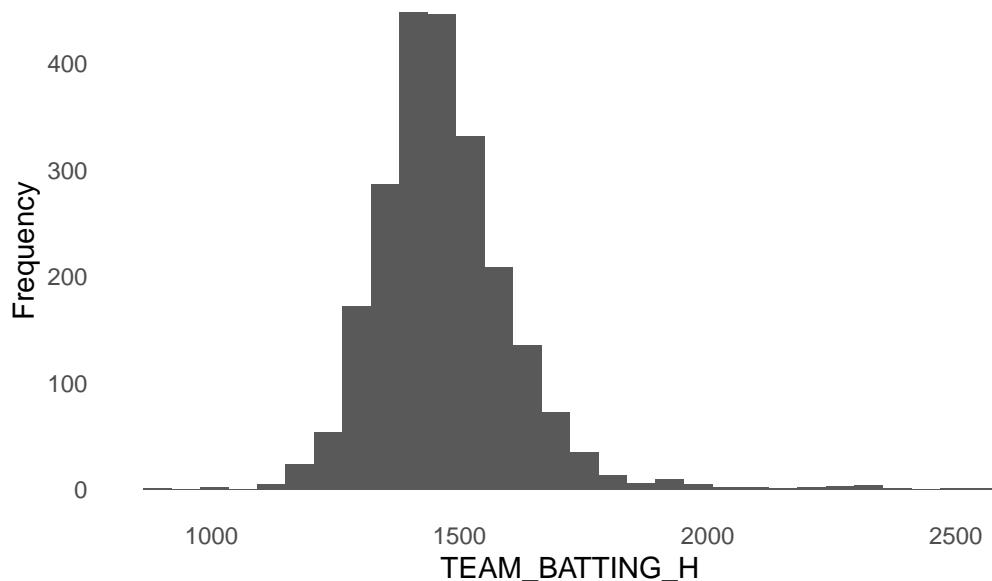
It makes sense that homeruns by batters (TEAM_BATTING_HR) and homeruns allowed (TEAM_PITCHING_HR) are correlated, as teams that hit many homeruns may also allow many homeruns, for example. We may need to avoid including these pairs of variables in the same model

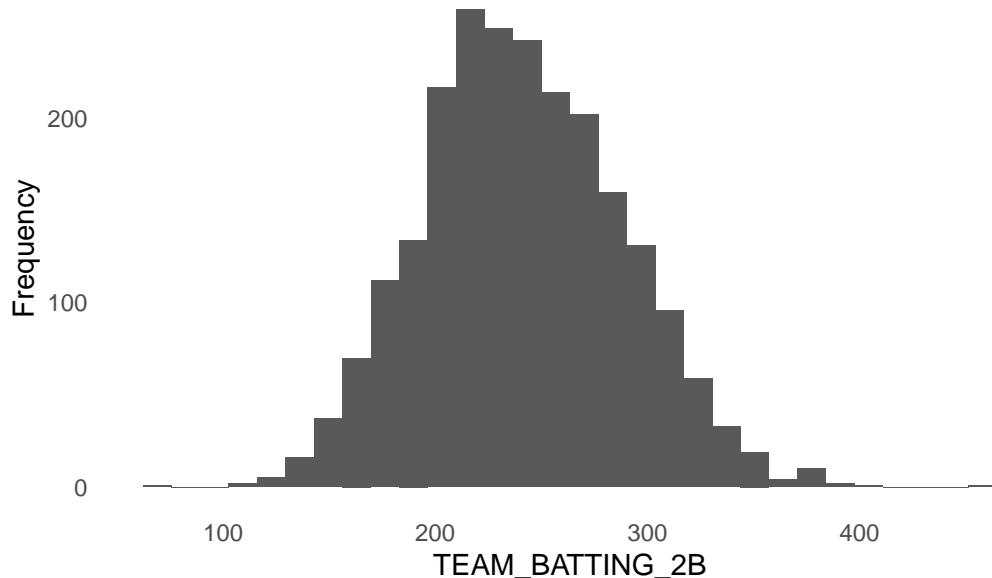
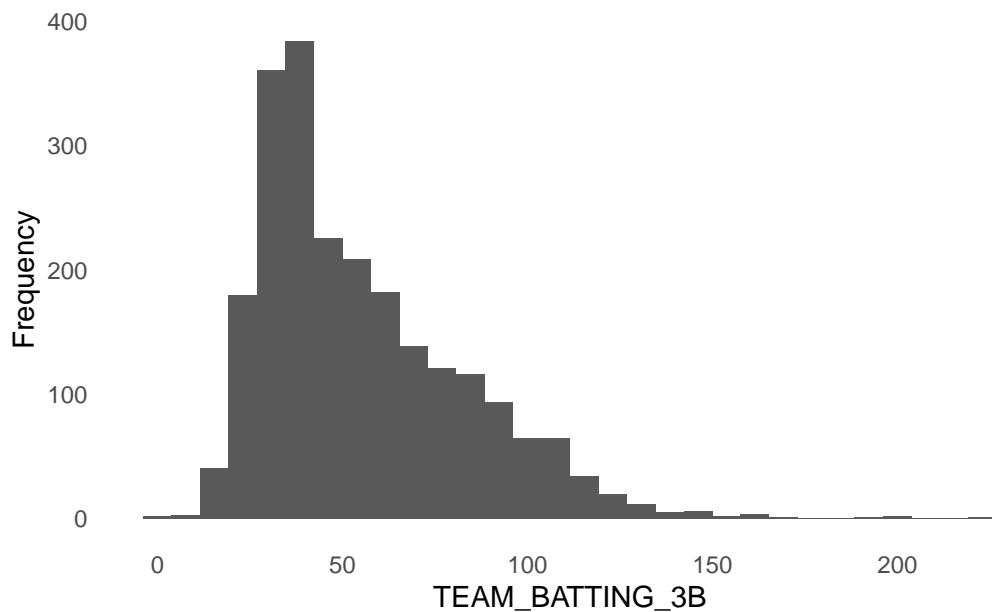
to prevent multicollinearity.

We also visualize the distribution of all variables.

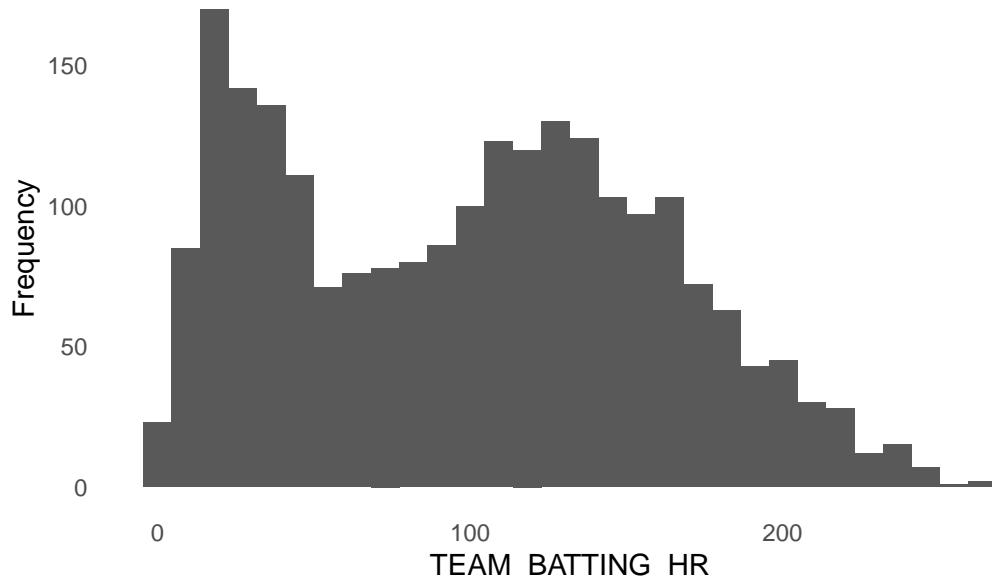
```
# Loop through all variables and create histograms
for (var in names(moneyball_train)) {
  p <- ggplot(moneyball_train, aes_string(x = var)) +
    geom_histogram() +
    labs(title = paste("Distribution of", var), x = var, y = "Frequency") +
    theme_minimal() +
    theme(panel.grid = element_blank())
  print(p)
}
```



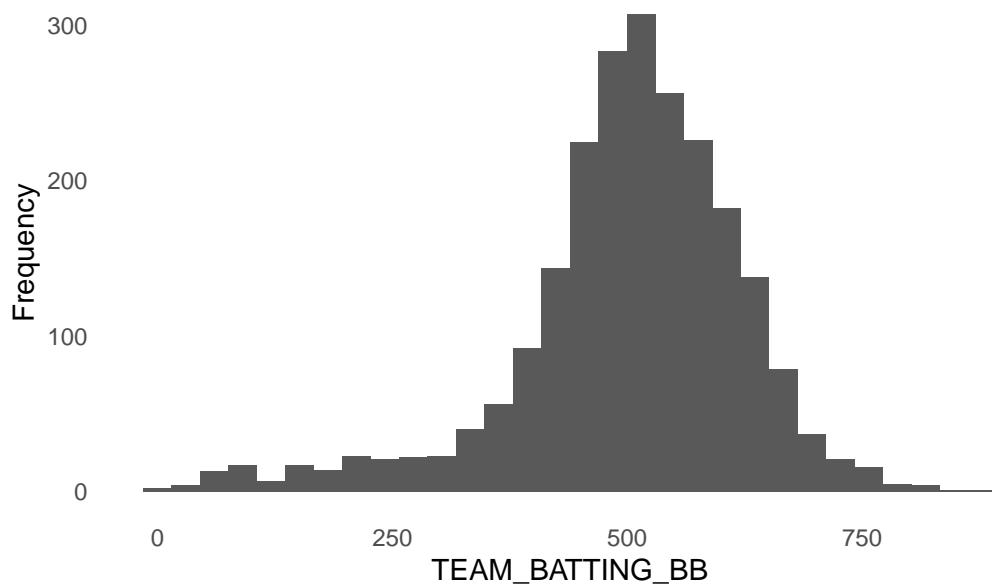
Distribution of TARGET_WINS**Distribution of TEAM_BATTING_H**

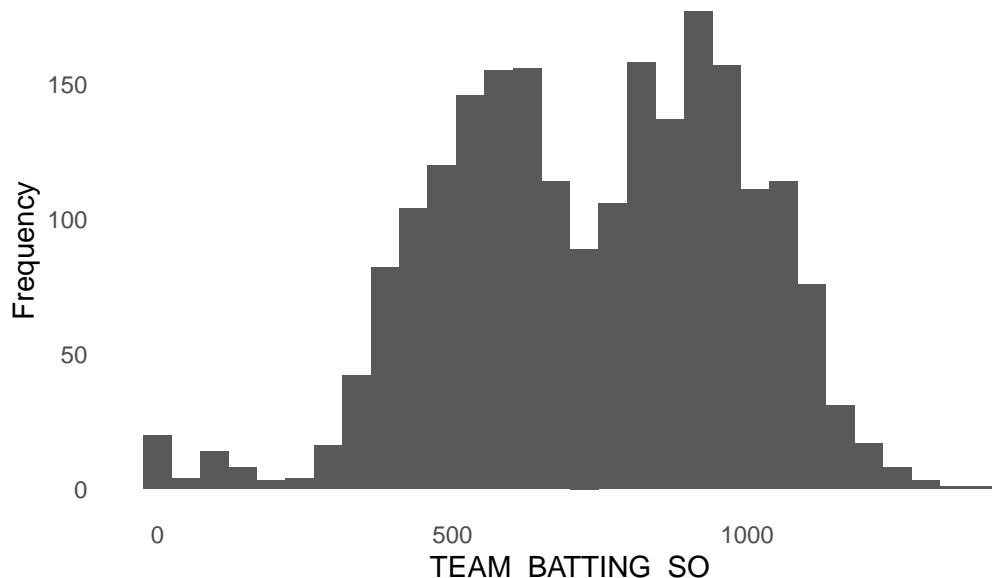
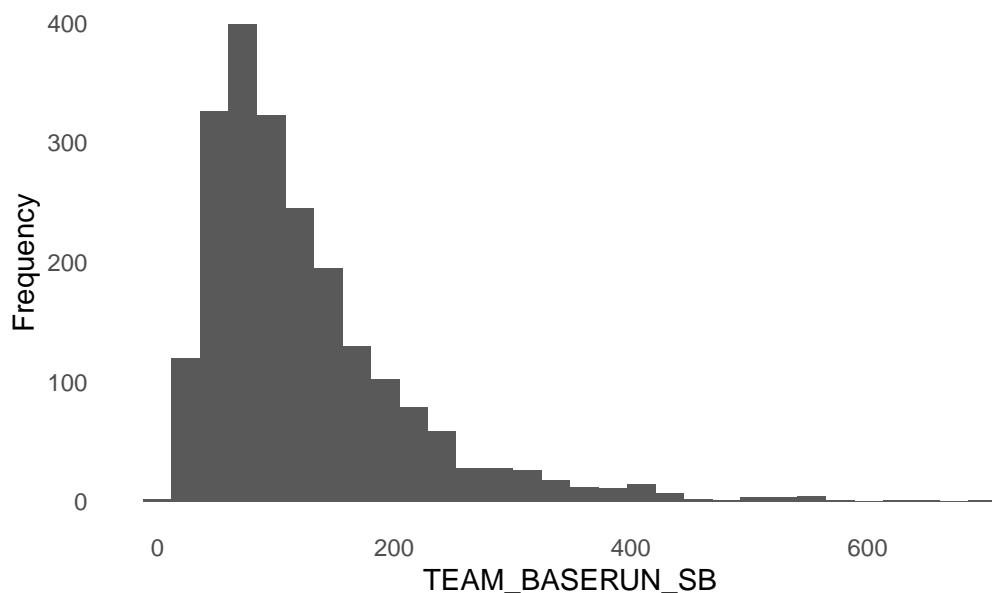
Distribution of TEAM_BATTING_2B**Distribution of TEAM_BATTING_3B**

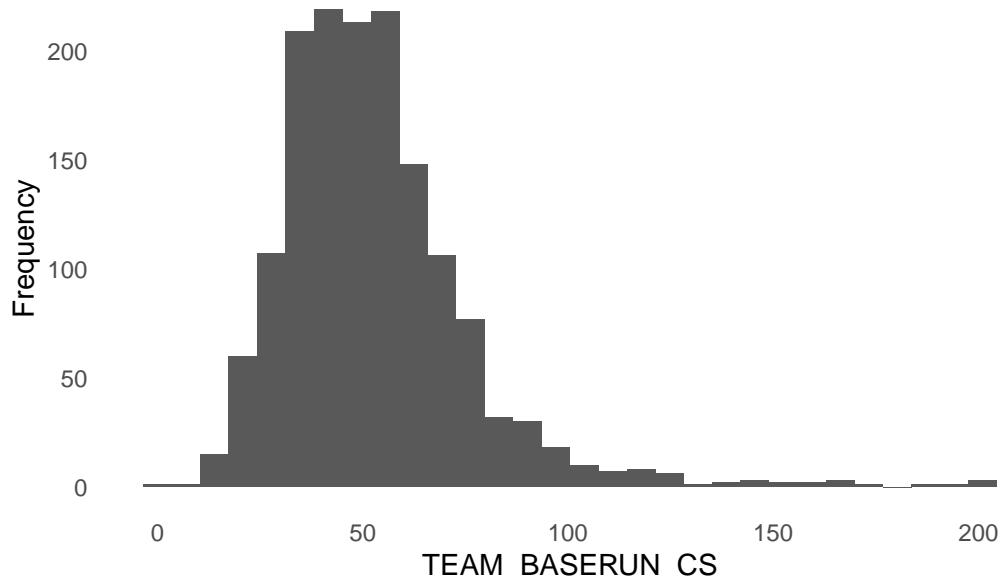
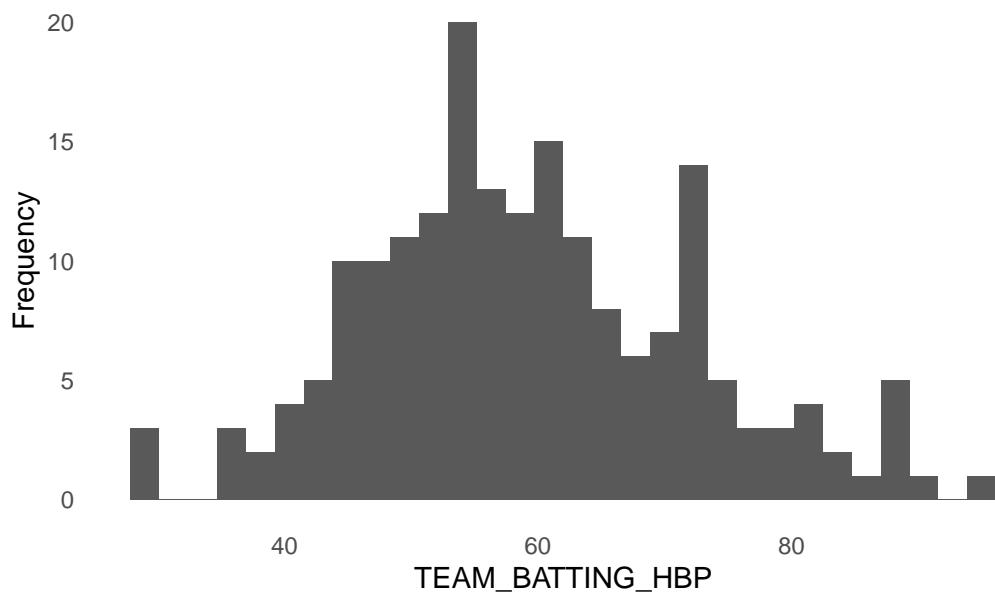
Distribution of TEAM_BATTING_HR

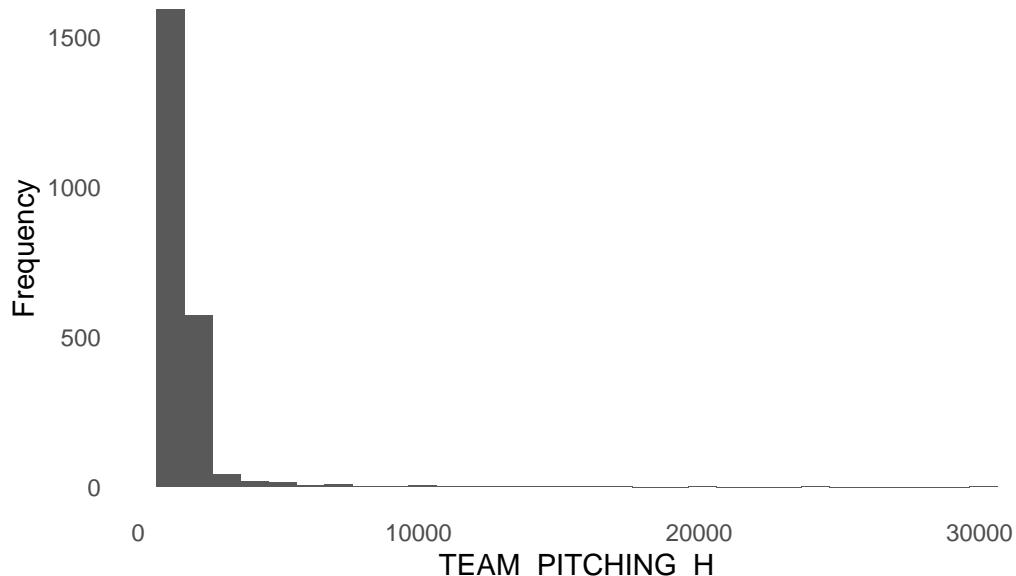
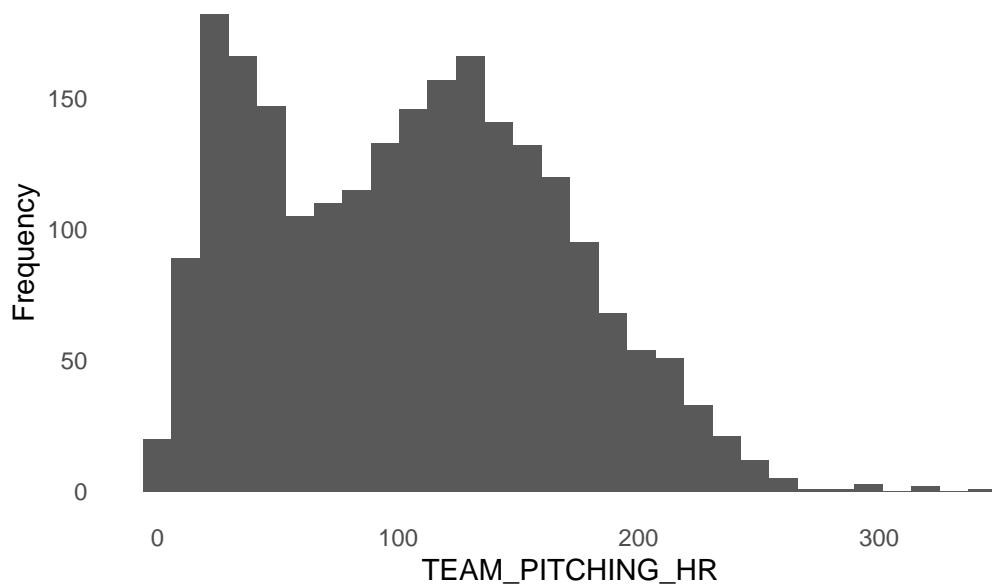


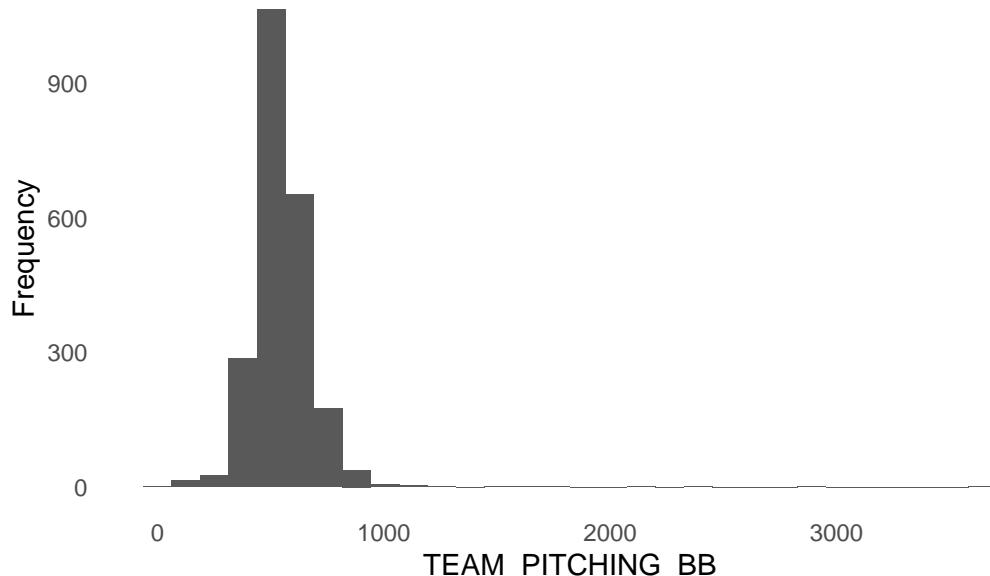
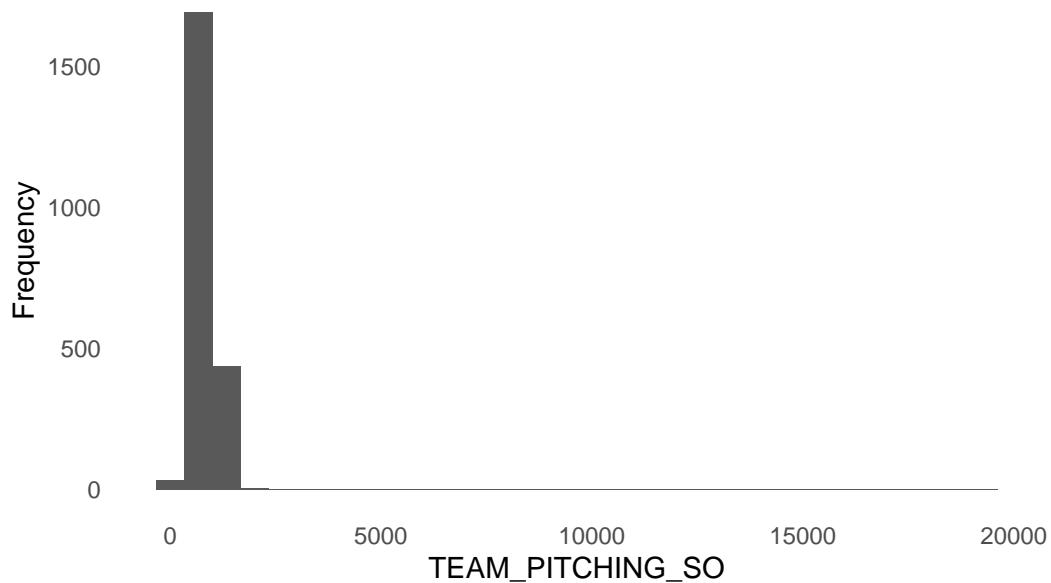
Distribution of TEAM_BATTING_BB

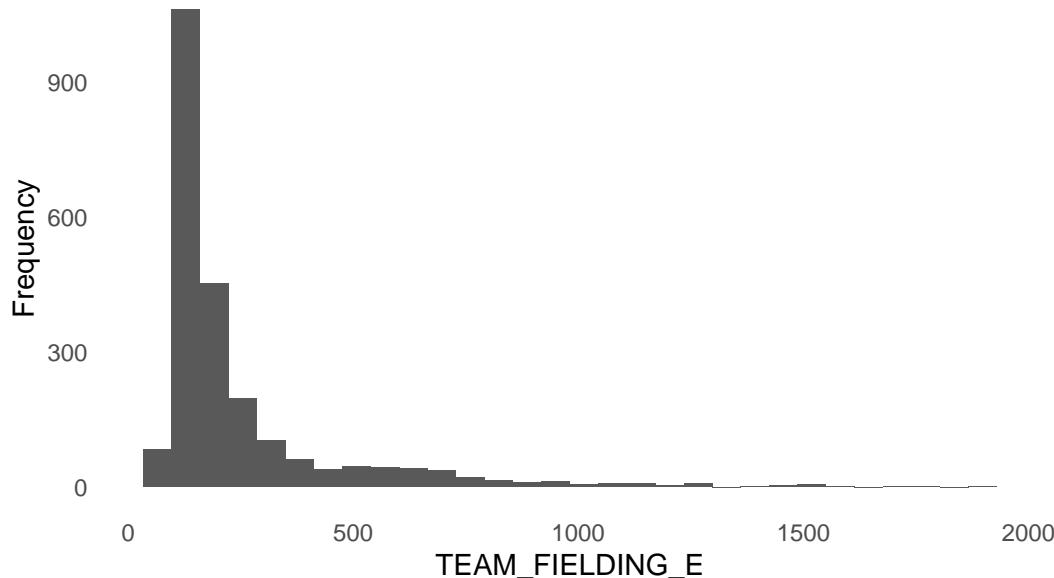
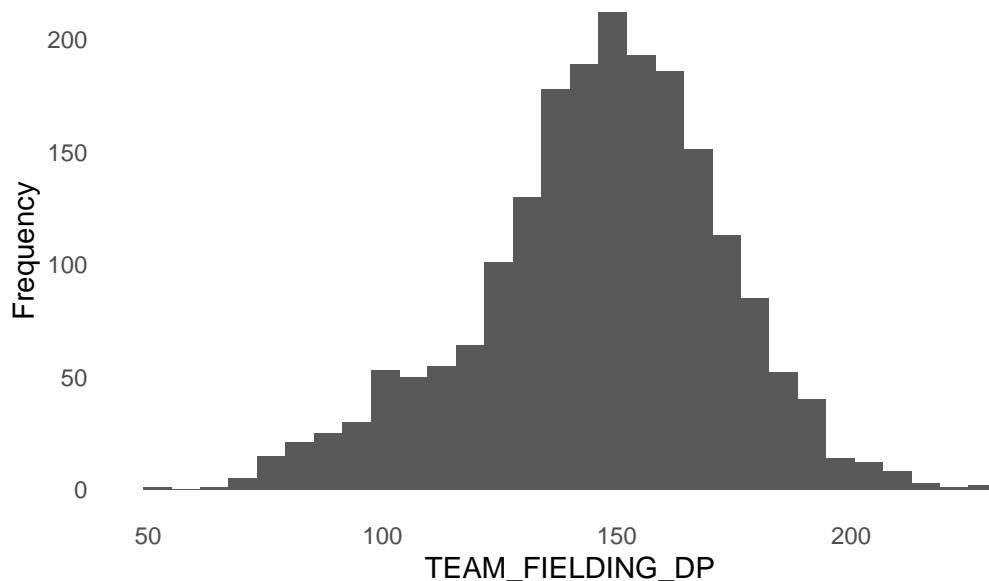


Distribution of TEAM_BATTING_SO**Distribution of TEAM_BASERUN_SB**

Distribution of TEAM_BASERUN_CS**Distribution of TEAM_BATTING_HBP**

Distribution of TEAM_PITCHING_H**Distribution of TEAM_PITCHING_HR**

Distribution of TEAM_PITCHING_BB**Distribution of TEAM_PITCHING_SO**

Distribution of TEAM_FIELDING_E**Distribution of TEAM_FIELDING_DP**

5.2 2. Data Preparation

Next, we will prepare the data for modeling.

First, we handle missing data. Since we saw in the EDA that TEAM_BATTING_HBP, TEAM_BASERUN_CS, and TEAM_FIELDING_DP were 91.6%, 33.9%, and 12.6% missing, respectively, and not very correlated with the target variable, so we will drop these variables. The remaining variables are less than 10% missing, so we will impute the missing values with the median of each variable. TEAM_BASERUN_SB

is the most correlated variable with the target variable that has missingness, so we will be sure to impute this variable rather than drop it.

The median imputation was picked based on the distributions of the variables with missingness: TEAM_BASERUN_SB and TEAM_PITCHING_SO have a long right tail so the median is a better imputation method than the mean. TEAM_BATTING_SO is bimodal, so the median is also a better imputation method than the mean for this variable. See the histograms above for the distributions of these variables.

```
moneyball_train_prepared <- moneyball_train |>  
  select(-TEAM_BATTING_HBP, -TEAM_BASERUN_CS, -TEAM_FIELDING_DP) |>  
  mutate(across(everything(), ~ ifelse(is.na(.), median(., na.rm = TRUE), .)))
```

5.3 3. Build Models

5.4 4. Select Models

Export assigned predictions (the number of wins for the team) for the evaluation data set for deliverable.