

Classification Metrics

Homework #2

Naomi Buell

Richie Rivera

Alexander Simon

Nana Frimpong

Said Naqwe

Invalid Date

```
library(tidyverse)
library(skimr)
library(corrplot)
library(ggplot2)
library(GGally)
library(janitor)
library(pracma)
```

1 Homework Solutions

1. Download the classification output data set (attached in Blackboard to the assignment).

```
classifications_output_data <- read.csv("classification-output-data (2).csv")

classifications_output_data |> skim()
```

Table 1: Data summary

Name	classifications_output_da...
Number of rows	181
Number of columns	11
Column type frequency:	
numeric	11
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
pregnant	0	1	3.86	3.24	0.00	1.00	3.00	6.00	15.00	
glucose	0	1	118.30	30.48	57.00	99.00	112.00	136.00	197.00	
diastolic	0	1	71.70	11.80	38.00	64.00	70.00	78.00	104.00	
skinfold	0	1	19.80	15.69	0.00	0.00	22.00	32.00	54.00	
insulin	0	1	63.77	88.73	0.00	0.00	0.00	105.00	543.00	
bmi	0	1	31.58	6.66	19.40	26.30	31.60	36.00	50.00	
pedigree	0	1	0.45	0.28	0.09	0.26	0.39	0.58	2.29	
age	0	1	33.31	11.18	21.00	24.00	30.00	41.00	67.00	
class	0	1	0.31	0.47	0.00	0.00	0.00	1.00	1.00	
scored.class	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
scored.probability	0	1	0.30	0.23	0.02	0.12	0.24	0.43	0.95	

2. The data set has three key columns we will use:

- *class*: the actual class for the observation
- *scored.class*: the predicted class for the observation (based on a threshold of 0.5)
- *scored.probability*: the predicted probability of success for the observation

Use the `table()` function to get the raw confusion matrix for this scored dataset. Make sure you understand the output. In particular, do the rows represent the actual or predicted class? The columns?

```
confusion_matrix <- table(
  classifications_output_data$class,
  classifications_output_data$scored.class
)
confusion_matrix
```

```
      0   1
0 119   5
1   30  27
```

```
tn <- confusion_matrix[1, 1]
fn <- confusion_matrix[2, 1]

# print
print(paste("True Negatives:", tn))
```

```
[1] "True Negatives: 119"
```

```
print(paste("False Negatives:", fn))
```

```
[1] "False Negatives: 30"
```

The rows represent the actual class (`class`) and the columns represent the predicted class (`scored.class`). For example, there are 119 observations where the actual class is 0 and the predicted class is also 0 (true negatives), and there are 30 observations where the actual class is 1 but the predicted class is 0 (false negatives).

3. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the accuracy of the predictions.

$$Accuracy = \frac{TP + TN}{TP + FP + TN + FN}$$

```
calculate_accuracy <- function(data, actual_col, predicted_col) {
  confusion_matrix <- table(data[[actual_col]], data[[predicted_col]])
  tp <- confusion_matrix[2, 2]
  tn <- confusion_matrix[1, 1]
  fp <- confusion_matrix[1, 2]
  fn <- confusion_matrix[2, 1]
  accuracy <- (tp + tn) / (tp + fp + tn + fn)
  return(accuracy)
}

# Test function
accuracy <- calculate_accuracy(classifications_output_data, "class", "scored.class")
print(paste("Accuracy:", accuracy))
```

```
[1] "Accuracy: 0.806629834254144"
```

4. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the classification error rate of the predictions.

$$ClassificationErrorRate = \frac{FP + FN}{TP + FP + TN + FN}$$

Verify that you get an accuracy and an error rate that sums to one.

```
calculate_error_rate <- function(data, actual_col, predicted_col) {
  confusion_matrix <- table(data[[actual_col]], data[[predicted_col]])
  tp <- confusion_matrix[2, 2]
  tn <- confusion_matrix[1, 1]
```

```

fp <- confusion_matrix[1, 2]
fn <- confusion_matrix[2, 1]
error_rate <- (fp + fn) / (tp + fp + tn + fn)
return(error_rate)
}

# Test function
error_rate <- calculate_error_rate(classifications_output_data, "class", "scored.class")
print(paste("Error Rate:", error_rate))

[1] "Error Rate: 0.193370165745856"

# Verify that accuracy and error rate sum to 1
print(paste("Sum of accuracy + error rate:", accuracy + error_rate))

[1] "Sum of accuracy + error rate: 1"

```

5. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the precision of the predictions.

$$Precision = \frac{TP}{TP + FP}$$

```

# Precision function

calculate_precision <- function(data, actual_col, predicted_col) {
  confusion_matrix <- table(data[[actual_col]], data[[predicted_col]])

  tp <- confusion_matrix[2, 2]
  tn <- confusion_matrix[1, 1]
  fp <- confusion_matrix[1, 2]
  fn <- confusion_matrix[2, 1]

  precision <- tp / (tp + fp)

  return(precision)
}

# Test function
precision <- calculate_precision(classifications_output_data, "class", "scored.class")
print(paste("Precision:", precision))

```

```
[1] "Precision: 0.84375"
```

6. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the sensitivity of the predictions. Sensitivity is also known as recall.

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

```
# Sensitivity function

calculate_sensitivity <- function(data, actual_col, predicted_col) {
  confusion_matrix <- table(data[[actual_col]], data[[predicted_col]])

  tp <- confusion_matrix[2, 2]
  tn <- confusion_matrix[1, 1]
  fp <- confusion_matrix[1, 2]
  fn <- confusion_matrix[2, 1]

  sensitivity <- tp / (tp + fn)

  return(sensitivity)
}

# Test function
sensitivity <- calculate_sensitivity(classifications_output_data, "class", "scored.class")
print(paste("Sensitivity:", sensitivity))
```

```
[1] "Sensitivity: 0.473684210526316"
```

7. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the specificity of the predictions.

$$\text{Specificity} = \frac{TN}{TN + FP}$$

```
calculate_specificity <- function(data, actual_col, predicted_col) {
  confusion_matrix <- table(data[[actual_col]], data[[predicted_col]])
  tp <- confusion_matrix[2, 2]
  tn <- confusion_matrix[1, 1]
  fp <- confusion_matrix[1, 2]
  fn <- confusion_matrix[2, 1]
  specificity <- tn / (tn + fp)
```

```

return(specificity)
}

# Test function
specificity <- calculate_specificity(classifications_output_data, "class", "scored.class")
print(paste("Specificity:", specificity))

```

```
[1] "Specificity: 0.959677419354839"
```

8. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the F1 score of the predictions.

$$F1Score = \frac{2 \times Precision \times Sensitivity}{Precision + Sensitivity}$$

```

# F1 score function

calculate_f1 <- function(data, actual_col, predicted_col) {

precision <- calculate_precision(data, actual_col, predicted_col)
sensitivity <- calculate_sensitivity(data, actual_col, predicted_col)

f1_score <- (2 * precision * sensitivity) / (precision + sensitivity)

return(f1_score)
}

# Test function
f1_score <- calculate_f1(classifications_output_data, "class", "scored.class")
print(paste("F1 Score:", f1_score))

```

```
[1] "F1 Score: 0.606741573033708"
```

9. Before we move on, let's consider a question that was asked: What are the bounds on the F1 score? Show that the F1 score will always be between 0 and 1. (Hint: If $0 < p < 1$ and $0 < s < 1$ then $2ps < p + s$.)

```

# F1 bounds demonstration

# Example values between 0 and 1
precision_test <- 0.7
sensitivity_test <- 0.8

```

```
f1_test <- (2 * precision_test * sensitivity_test) /
  (precision_test + sensitivity_test)

print(paste("Example F1 Score:", f1_test))
```

```
[1] "Example F1 Score: 0.7466666666666667"
```

10. Write a function that generates an ROC curve from a data set with a true classification column (class in our example) and a probability column (scored.probability in our example). Your function should return a list that includes the plot of the ROC curve and a vector that contains the calculated area under the curve (AUC). Note that I recommend using a sequence of thresholds ranging from 0 to 1 at 0.01 intervals.

```
# calc_auc() creates a ROC curve and calculates the area under the curve
# The inputs are a Boolean vector of known classifications ('labels') and
# a numeric vector of probability scores ('scores'). The function returns a
# list containing the ROC plot (ggplot object) and numeric AUC value
# Adapted from https://www.r-bloggers.com/2025/11/roc-curves-in-two-lines-of-code/

require(pracma)

calc_auc <- function(labels, scores){
  # Sort labels by score
  labels <- labels[order(scores, decreasing = TRUE)]
  # Calculate TPR and FPR
  rates_df <- data.frame(
    TPR = cumsum(labels) / sum(labels),
    # !labels is the negative class (0)
    FPR = cumsum(!labels) / sum(!labels))


  # Plot ROC curve
  p <- ggplot(rates_df, aes(x = FPR, y = TPR)) +
    geom_line() +
    ggtitle('ROC Curve')

  # Calculate AUC using trapezoidal approximation
  AUC <- pracma::trapz(rates_df$FPR, rates_df$TPR)

  return(list(roc_plot = p, auc_value = AUC))
}
```

```
# Call the AUROC function with the provided data
performance <- calc_auroc(classifications_output_data$class,
                          classifications_output_data$scored.probability)
```

```
# Show ROC plot
performance$roc_plot
```



```
# Show AUC value
sprintf('The area under the ROC curve is %.2f', performance$auc_value)
```

```
[1] "The area under the ROC curve is 0.85"
```

11. Use your created R functions and the provided classification output data set to produce all of the classification metrics discussed above.

```
accuracy <- calculate_accuracy(classifications_output_data, "class", "scored.class")
error_rate <- calculate_error_rate(classifications_output_data, "class", "scored.class")
precision <- calculate_precision(classifications_output_data, "class", "scored.class")
sensitivity <- calculate_sensitivity(classifications_output_data, "class", "scored.class")
specificity <- calculate_specificity(classifications_output_data, "class", "scored.class")
f1_score <- calculate_f1(classifications_output_data, "class", "scored.class")

# ROC and AUC
performance <- calc_auroc(classifications_output_data$class,
                          classifications_output_data$scored.probability)
```

```

# Summary table
metrics_table <- data.frame(
  Metric = c("Accuracy", "Error Rate", "Precision",
             "Sensitivity", "Specificity", "F1 Score", "AUC"),
  Value = round(c(accuracy, error_rate, precision,
                 sensitivity, specificity, f1_score,
                 performance$auc_value), 4)
)

# Table for Classification Metrics
knitr::kable(metrics_table,
              caption = "Classification Metrics Summary",
              col.names = c("Metric", "Value"),
              align = c("l", "r"))

```

Table 3: Classification Metrics Summary

Metric	Value
Accuracy	0.8066
Error Rate	0.1934
Precision	0.8438
Sensitivity	0.4737
Specificity	0.9597
F1 Score	0.6067
AUC	0.8503

12. Investigate the *caret* package. In particular, consider the functions *confusionMatrix*, *sensitivity*, and *specificity*. Apply the functions to the data set. How do the results compare with your own functions?

```

# Load caret package
library(caret)

```

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':

```
lift
```

```
# Create confusion matrix using caret
cm_caret <- confusionMatrix(as.factor(classifications_output_data$scored.class),
                             as.factor(classifications_output_data$class))

print(cm_caret)
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	119	30
1	5	27

```

      Accuracy : 0.8066
      95% CI : (0.7415, 0.8615)
No Information Rate : 0.6851
P-Value [Acc > NIR] : 0.0001712

```

```
      Kappa : 0.4916
```

```
McNemar's Test P-Value : 4.976e-05
```

```

      Sensitivity : 0.9597
      Specificity : 0.4737
      Pos Pred Value : 0.7987
      Neg Pred Value : 0.8438
      Prevalence : 0.6851
      Detection Rate : 0.6575
      Detection Prevalence : 0.8232
      Balanced Accuracy : 0.7167

```

```
'Positive' Class : 0
```

```
# Extract metrics from caret
caret_accuracy <- cm_caret$overall['Accuracy']
caret_sensitivity <- cm_caret$byClass['Sensitivity']
caret_specificity <- cm_caret$byClass['Specificity']

print(paste("Caret Accuracy:", caret_accuracy))
```

```
[1] "Caret Accuracy: 0.806629834254144"
print(paste("Caret Sensitivity:", caret_sensitivity))

[1] "Caret Sensitivity: 0.959677419354839"
print(paste("Caret Specificity:", caret_specificity))

[1] "Caret Specificity: 0.473684210526316"
# Compare with our functions
print("Comparison:")

[1] "Comparison:"
print(paste("Accuracy - Ours:", accuracy, "Caret:", caret_accuracy))

[1] "Accuracy - Ours: 0.806629834254144 Caret: 0.806629834254144"
print(paste("Sensitivity - Ours:", sensitivity, "Caret:", caret_sensitivity))

[1] "Sensitivity - Ours: 0.473684210526316 Caret: 0.959677419354839"
print(paste("Specificity - Ours:", specificity, "Caret:", caret_specificity))

[1] "Specificity - Ours: 0.959677419354839 Caret: 0.473684210526316"
```

Both methods calculated an identical accuracy of 0.8066 , but the sensitivity and specificity values were perfectly swapped. This discrepancy occurred because our custom functions evaluated '1' as the target condition, whereas the caret package defaulted to treating '0' as the positive class.

13. Investigate the *pROC* package. Use it to generate an ROC curve for the data set. How do the results compare with your own functions?

```
# Load pROC package
library(pROC)
```

Type 'citation("pROC")' for a citation.

Attaching package: 'pROC'

The following objects are masked from 'package:stats':

```
cov, smooth, var
```


```
# Generate ROC curve using pROC
roc_obj <- roc(classifications_output_data$class, classifications_output_data$scored.proba
```

```
Setting levels: control = 0, case = 1
```

```
Setting direction: controls < cases
```

```
# Plot the ROC curve
```

```
plot(roc_obj, main = "ROC Curve using pROC")
```



```
# Get AUC
```

```
auc_proc <- auc(roc_obj)
```

```
print(paste("AUC from pROC:", auc_proc))
```

```
[1] "AUC from pROC: 0.850311262026033"
```

```
# Compare with our custom function (assuming roc_result$auc is available from section 10)
```

```
print("Comparison:")
```

```
[1] "Comparison:"
```

```
print(paste("Custom AUC:", roc_obj$auc, "pROC AUC:", auc_proc))
```

```
[1] "Custom AUC: 0.850311262026033 pROC AUC: 0.850311262026033"
```

The results from the pROC package are mathematically identical to our custom functions, with both yielding an AUC of 0.8503. The primary difference are in the visualization defaults. Our custom ggplot function plots True Positive Rate vs. False Positive Rate (0 to 1). In contrast, pROC's base plot graphs Sensitivity vs. Specificity, drawing the Specificity x-axis in reverse order (1 to 0) to maintain the standard ROC curve shape.